

Lecture 2 – Distributed Filesystems

922EU3870 – Cloud Computing and
Mobile Platforms, Autumn 2009

2009/9/21

Ping Yeh (葉平), Google, Inc.

Outline

- Get to know the numbers
- Filesystems overview
- Distributed file systems
 - Basic (example: NFS)
 - Shared storage (example: Global FS)
 - Wide-area (example: AFS)
 - Fault-tolerant (example: Coda)
 - Parallel (example: Lustre)
 - Fault-tolerant and Parallel (example: dCache)
- The Google File System
- Homework



Numbers real world engineers should know

L1 cache reference	0.5 ns
Branch mispredict	5 ns
L2 cache reference	7 ns
Mutex lock/unlock	100 ns
Main memory reference	100 ns
Compress 1 KB with Zippy	10,000 ns
Send 2 KB through 1 Gbps network	20,000 ns
Read 1 MB sequentially from memory	250,000 ns
Round trip within the same data center	500,000 ns
Disk seek	10,000,000 ns
Read 1 MB sequentially from network	10,000,000 ns
Read 1 MB sequentially from disk	30,000,000 ns
Round trip between California and Netherlands	150,000,000 ns



The Joys of Real Hardware

Typical first year for a new cluster:

- ~0.5 **overheating** (power down most machines in <5 mins, ~1-2 days to recover)
- ~1 **PDU failure** (~500-1000 machines suddenly disappear, ~6 hours to come back)
- ~1 **rack-move** (plenty of warning, ~500-1000 machines powered down, ~6 hours)
- ~1 **network rewiring** (rolling ~5% of machines down over 2-day span)
- ~20 **rack failures** (40-80 machines instantly disappear, 1-6 hours to get back)
- ~5 **racks go wonky** (40-80 machines see 50% packetloss)
- ~8 **network maintenances** (4 might cause ~30-minute random connectivity losses)
- ~12 **router reloads** (takes out DNS and external vips for a couple minutes)
- ~3 **router failures** (have to immediately pull traffic for an hour)
- ~dozens of minor **30-second blips for dns**
- ~1000 **individual machine failures**
- ~thousands of **hard drive failures**

slow disks, bad memory, misconfigured machines, flaky machines, etc.

File Systems Overview

- System that permanently stores data
- Usually layered on top of a lower-level physical storage medium
- Divided into logical units called “files”
 - Addressable by a filename (“foo.txt”)
- Files are often organized into directories
 - Usually supports hierarchical nesting (directories)
 - A path is the expression that joins directories and filename to form a unique “full name” for a file.
- Directories may further belong to a volume
- The set of valid paths form the *namespace* of the file system.



What Gets Stored

- User data itself is the bulk of the file system's contents
- Also includes meta-data on a volume-wide and per-file basis:

Volume-wide:

Available space

Formatting info

character set

...

Per-file:

name

owner

modification date

physical layout...



High-Level Organization

- Files are typically organized in a “tree” structure made of nested directories
- One directory acts as the “root”
- “links” (symlinks, shortcuts, etc) provide simple means of providing multiple access paths to one file
- Other file systems can be “mounted” and dropped in as sub-hierarchies (other drives, network shares)
- Typical operations on a file: create, delete, rename, open, close, read, write, append.
 - also lock for multi-user systems.



Low-Level Organization (1/2)

- File data and meta-data stored separately
- File descriptors + meta-data stored in inodes (Un*x)
 - Large tree or table at designated location on disk
 - Tells how to look up file contents
- Meta-data may be replicated to increase system reliability



Low-Level Organization (2/2)

- “Standard” read-write medium is a hard drive (other media: CDROM, tape, ...)
- Viewed as a sequential array of blocks
- Usually address ~1 KB chunk at a time
- Tree structure is “flattened” into blocks
- Overlapping writes/deletes can cause fragmentation: files are often not stored with a linear layout
 - inodes store all block numbers related to file



Fragmentation

A	B	C	(free space)
---	---	---	--------------

A	B	C	A	(free space)
---	---	---	---	--------------

A	(free space)	C	A	(free space)
---	--------------	---	---	--------------

A	D	C	A	D	(free)
---	---	---	---	---	--------




Filesystem Design Considerations

- Namespace: physical, logical
- Consistency: what to do when more than one user reads/writes on the same file?
- Security: who can do what to a file? Authentication/ACL
- Reliability: can files not be damaged at power outage or other hardware failures?




Local Filesystems on Unix-like Systems

- Many different designs
- Namespace: root directory “/”, followed by directories and files.
- Consistency: “sequential consistency”, newly written data are immediately visible to open reads (if...)
- Security:
 - uid/gid, mode of files
 - kerberos: tickets
- Reliability: superblocks, journaling, snapshot
 - more reliable filesystem on top of existing filesystem: RAID




Namespace



- Physical mapping: a directory and all of its subdirectories are stored on the same physical media.
 - /mnt/cdrom
 - /mnt/disk1, /mnt/disk2, ... when you have multiple disks
- Logical volume: a logical namespace that can contain multiple physical media or a partition of a physical media
 - still mounted like /mnt/vol1
 - dynamical resizing by adding/removing disks without reboot
 - splitting/merging volumes as long as no data spans the split




Journaling



- Changes to the filesystem is logged in a *journal* before it is committed.
 - useful if an supposedly atomic action needs two or more writes, e.g., appending to a file (update metadata + allocate space + write the data).
 - can play back a journal to recover data quickly in case of hardware failure. (old-fashioned “scan the volume” fsck anyone?)
- What to log?
 - changes to file content: heavy overhead
 - changes to metadata: fast, but data corruption may occur
- Implementations: xfs3, ReiserFS, IBM's JFS, etc.




RAID



- Redundant Array of Inexpensive Disks
- Different RAID levels
 - RAID 0 (striped disks): data is distributed among disk for performance.
 - RAID 1 (mirror): data is mirrored on another disk 1:1 for reliability.
 - RAID 5 (striped disks with distributed parity): similar to RAID 0, but with a parity bit for data recovery in case one disk is lost. The parity bit is rotated among disks to maximize throughput.
 - RAID 6 (striped disks with dual distributed parity): similar to RAID 5 but with 2 parity bits. Can lose 2 disks without losing data.
 - RAID 10 (1+0, striped mirrors)



Snapshot



- A snapshot = a copy of a set of files and directories at a point in time.
 - read-only snapshots, read-write snapshots
 - usually done by the filesystem itself, sometimes by LVMs
 - backing up data can be done on a read-only snapshot without worrying about consistency
- Copy-on-write is a simple and fast way to create snapshots
 - current data is the snapshot.
 - a request to write to a file creates a new copy, and work from there afterwards.
- Implementation: UFS, Sun's ZFS, etc.



- Should the file system be faster or more reliable?
- But faster at what: Large files? Small files? Lots of reading? Frequent writers, occasional readers?
- Block size
 - Smaller block size reduces amount of wasted space
 - Larger block size increases speed of sequential reads (may not help random access)



Distributed File Systems



Distributed Filesystems

- Support access to files on remote servers
- Must support concurrency
 - Make varying guarantees about locking, who “wins” with concurrent writes, etc...
 - Must gracefully handle dropped connections
- Can offer support for replication and local caching
- Different implementations sit in different places on complexity/feature scale



DFS is useful when...

- Multiple users want to share files
- Users move from one computer to another
- Users want a uniform namespace for shared files on every computer
- Users want a centralized storage system – backup and management
- Note that a “user” of a DFS may actually be a “program”




Features of Distributed File Systems

- Different systems have different designs and behaviors on the following features.
 - Interface: file system, block I/O, custom made
 - Security: various authentication/authorization schemes.
 - Reliability (fault-tolerance): can continue to function when some hardware fail (disks, nodes, power, etc).
 - Namespace (virtualization): can provide logical namespace that can span across physical boundaries.
 - Consistency: all clients get the same data all the time. It is related to locking, caching, and synchronization.
 - Parallel: multiple clients can have access to multiple disks at the same time.
 - Scope: local area network vs. wide area network.




NFS

- First developed in 1980s by Sun
- Most widely known distributed filesystem.
- Presented with standard POSIX FS interface
- Network drives are mounted into local directory hierarchy
 - Your home directory at CINC is probably NFS-driven
 - Type 'mount' some time at the prompt if curious




NFS Protocol



- Initially completely stateless
 - Operated over UDP; did not use TCP streams
 - File locking, etc, implemented in higher-level protocols
 - Implication: All client requests have enough info to complete op
 - example: Client specifies offset in file to write to
 - one advantage: Server state does not grow with more clients
- Modern implementations use TCP/IP & stateful protocols
- RFCs:
 - RFC 1094 for v2 (3/1989)
 - RFC 1813 for v3 (6/1995)
 - RFC 3530 for v4 (4/2003)




NFS Overview



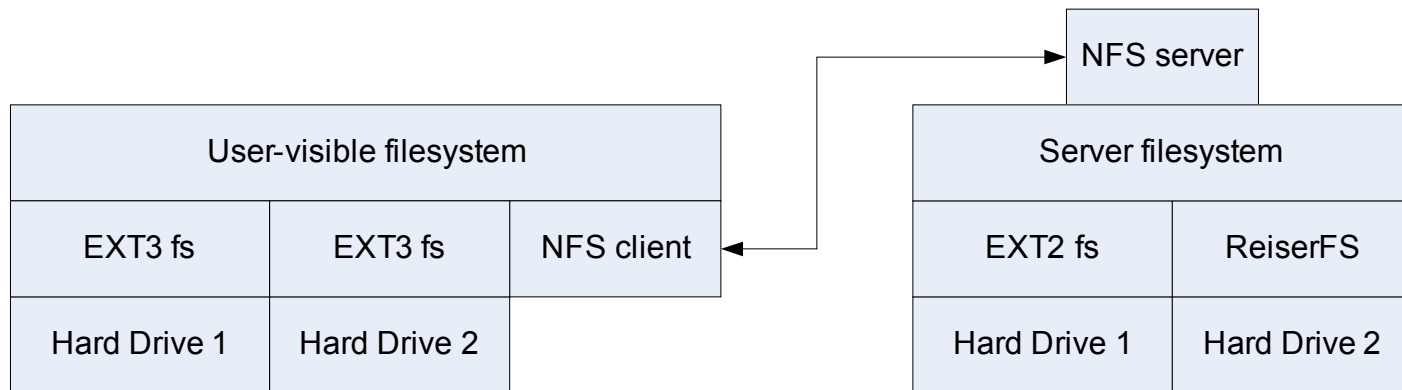
- Remote Procedure Calls (RPC) for communication between client and server
- Client Implementation
 - Provides transparent access to NFS file system
 - UNIX contains Virtual File system layer (VFS)
 - Vnode: interface for procedures on an individual file
 - Translates vnode operations to NFS RPCs
- Server Implementation
 - Stateless: Must not have anything only in memory
 - Implication: All modified data written to stable storage before return control to client
 - Servers often add NVRAM to improve performance




Server-side Implementation



- NFS defines a virtual file system
 - Does not actually manage local disk layout on server
- Server instantiates NFS volume on top of local file system
 - Local hard drives managed by concrete file systems (ext*, ReiserFS, ...)
 - Export local disks to clients




NFS Locking



- NFS v4 supports stateful locking of files
 - Clients inform server of intent to lock
 - Server can notify clients of outstanding lock requests
 - Locking is lease-based: clients must continually renew locks before a timeout
 - Loss of contact with server abandons locks




NFS Client Caching



- NFS Clients are allowed to cache copies of remote files for subsequent accesses
- Supports *close-to-open* cache consistency
 - When client A closes a file, its contents are synchronized with the master, and timestamp is changed
 - When client B opens the file, it checks that local timestamp agrees with server timestamp. If not, it discards local copy.
 - Concurrent reader/writers must use flags to disable caching




Cache Consistency



- Problem: Consistency across multiple copies (server and multiple clients)
- How to keep data consistent between client and server?
 - If file is changed on server, will client see update?
 - Determining factor: Read policy on clients
- How to keep data consistent across clients?
 - If write file on client A and read on client B, will B see update?
 - Determining factor: Write and read policy on clients



NFS: Summary




- Very popular
- Full POSIX interface means it “drops in” very easily.
- No cluster-wide uniform namespace: each client decides how to name a volume mounted from an NFS server.
- No location transparency: data movement means remount
- NFS volume managed by single server
 - Higher load on central server, bottleneck on scalability
 - Simplifies coherency protocols
- Challenges: Fault tolerance, scalable performance, and consistency




AFS (The Andrew File System)

- Developed at Carnegie Mellon University since 1983
- Strong security: Kerberos authentication
 - richer set of access control bits than UNIX: separate “administer”/“delete” bits, application-specific bits.
- Higher scalability compared with distributed file systems at the time. Supports 50,000+ clients at enterprise level.
- Global namespace: /afs/cmu.edu/vol1/..., WAN scope
- Location transparency: moving files just works.
- Heavily use client side caching.
- Support read-only replication.




Morgan Stanley's comparison: AFS vs. NFS



- Better client/server ratio
 - NFS (circa 1993) topped out at 25:1
 - AFS: 100s
- Robust volume replication
 - NFS servers go down, and take their clients with them
 - AFS servers go down, nobody notices (OK, RO data only)
- WAN file sharing
 - NFS couldn't do it reliably
 - AFS worked like a charm
- Security was never an issue (kerberos4)


Local Caching



- File reads/writes operate on locally cached copy
- The modified local copy is sent back to server when file is closed
- Open local copies are notified of external updates through callbacks
 - but not updated
- Trade-offs:
 - Shared database files do not work well on this system
 - Does not support write-through to shared medium




Replication



- AFS allows read-only copies of filesystem volumes
- Copies are guaranteed to be atomic checkpoints of entire FS at time of read-only copy generation (“snapshot”)
- Modifying data requires access to the sole r/w volume
 - Changes do not propagate to existing read-only copies



AFS: Summary



- Not quite POSIX
 - Stronger security/permissions
 - No file write-through
- High availability through replicas and local caching
- Scalability by reducing load on servers
- Not appropriate for all file types

ACM Trans. on Computer Systems, 6(1):51-81, Feb 1988.




Global File System




Lustre Overview

- Developed by Peter Braam at Carnegie Mellon since 1999
 - then Cluster File System in 2001, acquired by Sun in 2007
 - Goal: removing bottlenecks to achieve scalability
- Object-based: separate metadata and file data
 - metadata server(s): MDS (namespace, ACL, attributes)
 - object storage server(s): OSS
 - client read/write directly with OSS
- Consistency: Lustre distributed lock manager
- Performance: data can be striped
- POSIX interface



Key Design Issue : Scalability



- I/O throughput
 - How to avoid bottlenecks
- Metadata scalability
 - How can 10,000's of nodes work on files in same folder
- Cluster Recovery
 - If sth fails, how can transparent recovery happen
- Management
 - Adding, removing, replacing, systems; data migration & backup



The Lustre Architecture






Figure 1. Lustre architecture for clusters


Metadata Service (MDS)



- All access to the file is governed by MDS which will directly or indirectly authorize access.
- To control namespace and manage inodes
- Load balanced cluster service for the scalability (a well balanced API, a stackable framework for logical MDS, replicated MDS)
- Journaled batched metadata updates




Object Storage Targets (OST)



- Keep file data objects
- File I/O service: Access to the objects
- The block allocation for data obj., leading distributed and scalability
- OST s/w modules
 - OBD server, Lock server
 - Obj. storage driver, OBD filter
 - Portal API



Distributed Lock Manager



- For generic and rich lock service
- Lock resources: resource database
 - Organize resources in trees
- High performance
 - node that acquires resource manages tree
- Intent locking: utilization dependent locking modes
 - write-back lock for low utilization files
 - no lock for high contention files, write to OSS by DLM



Metadata

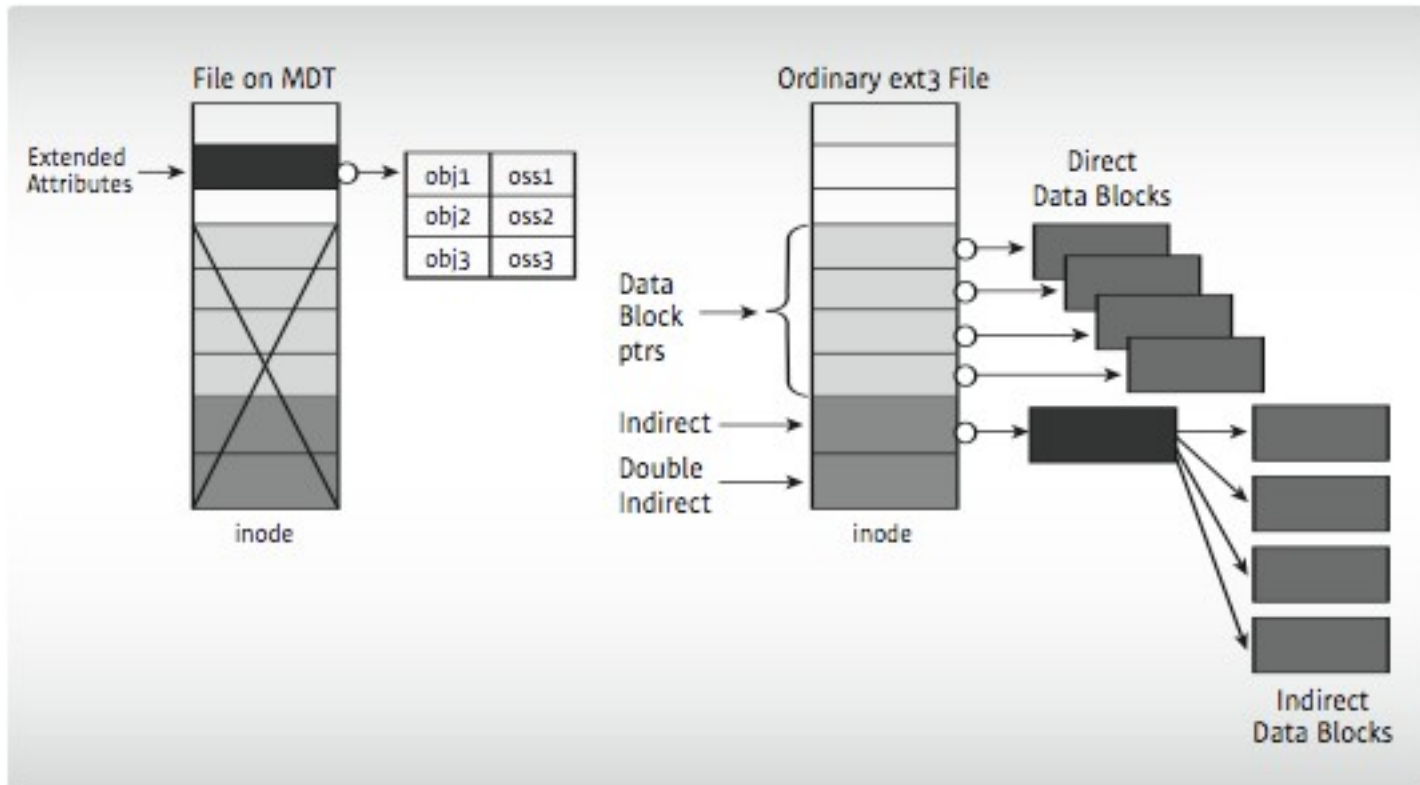



Figure 5. MDS inodes point to objects; ext3 inodes point to data



File I/O

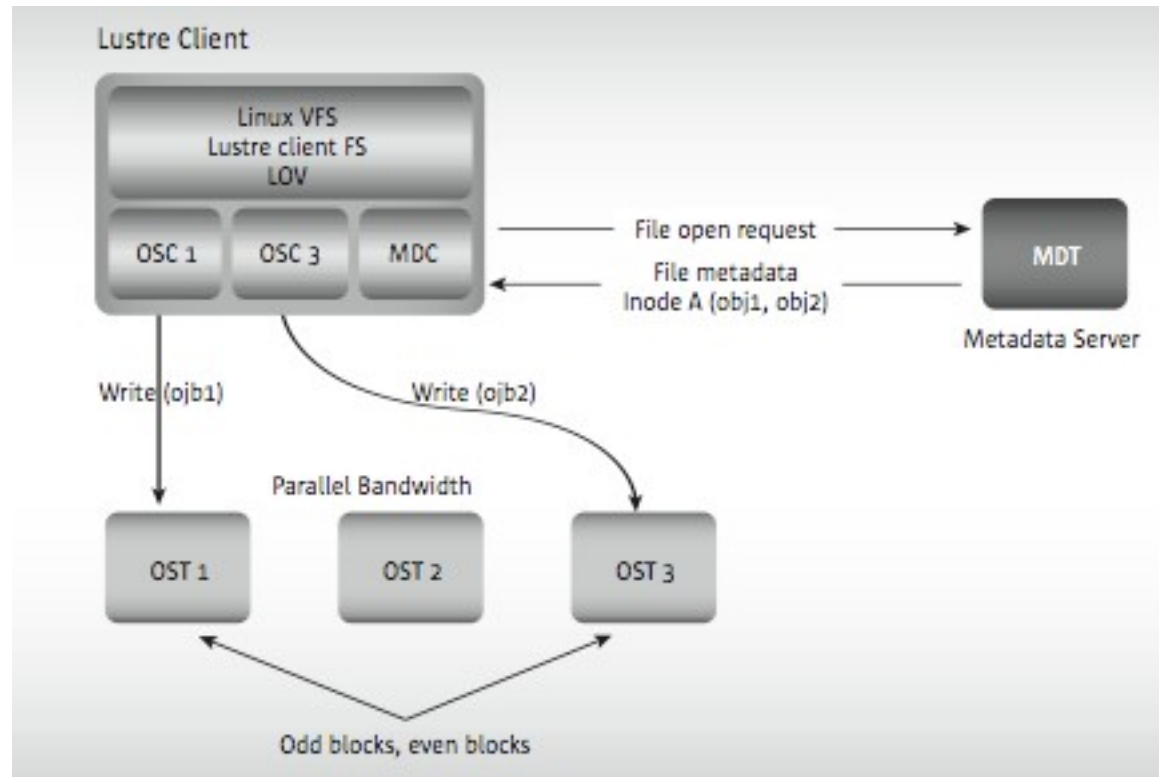



Figure 6. File open and file I/O in the Lustre file system



Striping

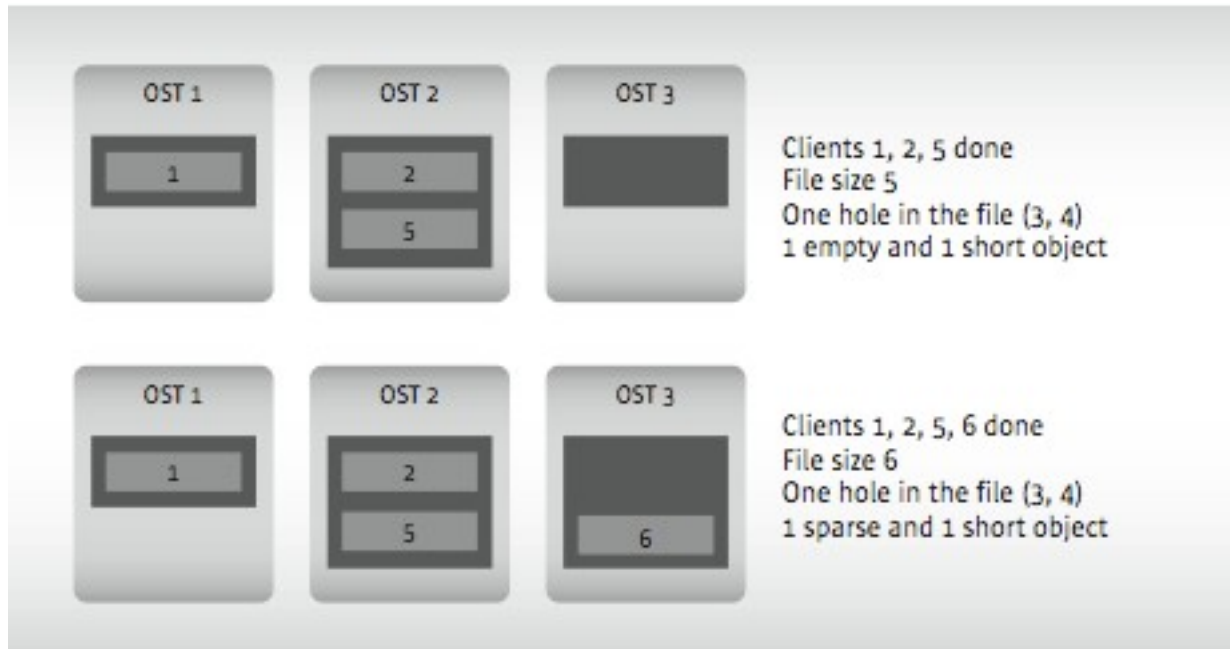




Figure 8. A compute application rendering a movie file with three objects



Lustre Summary



- Good scalability and high performance
 - Lawrence Livermore National Lab BlueGene/L has 200,000 clients processes access 2.5PB Lustre fs through 1024 IO nodes over 1Gbit Ethernet network at 35 GB/sec.
- Reliability
 - Journaled metadata in metadata cluster
 - OSTs are responsible for their data
- Consistency



Other file systems worth mentioning

- Coda: post-AFS, pre-Lustre, from CMU
- Global File System: shared disk file systems
- PVFS by IBM
- ZFS by Sun

